



Boat Listing Case Study

Ademola Olokun

Boat Listing Insights

December 2022

Introduction

- Users can advertise their used boats for sale on Nearly New Nautical. Listings require users to provide a range of information about their boats. The more views a boat receives, the more traffic the website receives, and the more potential customers there are.
- In order to boost traffic to the website, the product manager wants to eliminate boats that do not receive many views from being listed by building a model that:
 - Predicts the number of views a listing will receive based on boat features
 - On average, predicts number of views that are only 50% off of the true number of views a listing would receive



A person is working at a desk in a modern office. A large monitor displays a data dashboard with various charts and tables. A laptop is open to the left, and a coffee cup sits on the desk. The person's hands are visible, one holding a smartphone and the other near the laptop. The scene is brightly lit, suggesting a sunny day.

Methodology


- The dataset consists of 9,888 boat listings with 10 features, including the target label, Number of Views.
- I am choosing to use CatBoost – a gradient boosting library, because we have more categorical features with moderate cardinality and CatBoost is better equipped to handle categorical features without negatively affecting the performance of the model. A Linear Regression model will serve as the baseline model.
- For the evaluation metric, I am choosing Mean Absolute Percentage Error to evaluate the model, which measures the prediction accuracy as a loss function of our regression model.
- The MAPE of our best performing predictive model is 0.5 (50%).

Outcome

	RMSE	MAPE	R2
model_name			
LinearRegression	129.700000	0.800000	0.170000
CatBoostRegressor	111.460000	0.500000	0.390000

The **best performing model** for the business case is the **CatBoostRegressor model**, with a MAPE value of 0.5

Outcome




Our KPI, Mean Absolute Percentage Error for the Linear Regression model, is 0.8 (80%).

This means, on average, our predicted views will be 80% off the true number of views.

As seen from the chart, a subset of the data shows that predicted views and true views differ greatly.

Outcome




Our KPI, Mean Absolute Percentage Error for the CatBoostRegressor model, is 0.5 (50%).

This means, on average, our predicted views will be 50% off the true number of views.

As seen from the chart, a subset of the data shows that predicted views and true views are mostly overlapping.

Outcome



From the chart, we can see the top features that impact our CatBoostRegressor model greatly:

- **Location**
- **Boat Type**
- **Price**
- Manufacturer
- Material
- Age
- Boat Condition
- Length
- Width
- Engine Type

Recommendations

To help increase traffic to the website, we can deploy this model to filter out boat listings that do not meet a threshold of number of views.

To better improve the model, I would recommend changing the input for the Location field from country to regions, to reduce the number of unique values, and potentially improving the performance of the model.

I would also recommend that the data fields are validated to minimize the occurrence of missing values, improving data analysis and the performance of the model.



A blurred background image showing several people in a meeting or office setting. The focus is on the silhouettes and out-of-focus faces of individuals, with warm lighting from the right side.

Thank you
